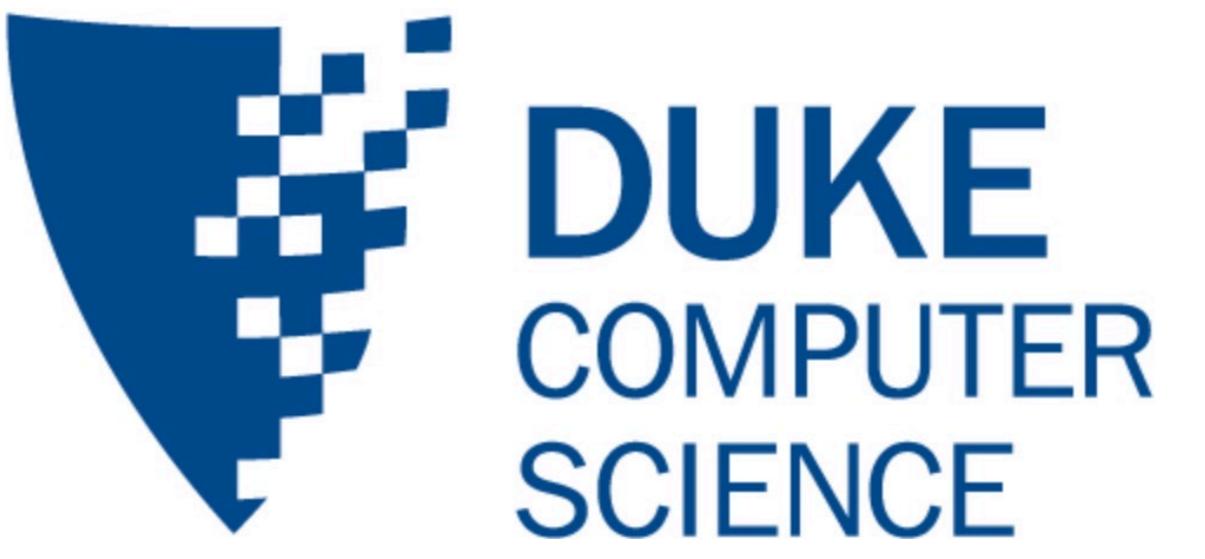




Optical Flow Training under Limited Label Budget via Active Learning



Shuai Yuan, Xian Sun, Hannah Kim, Shuzhi Yu, and Carlo Tomasi
Department of Computer Science, Duke University, USA

Introduction

“Accuracy-annotation cost” trade-off

- Optical flow: especially **hard** to label
- **Supervised**: better accuracy
- Unsupervised**: lower annotation cost
- Semi-supervised**: balancing both
- Draw “**Error-label ratio**” curves to visualize **label efficiency**
- Even a few labels help significantly

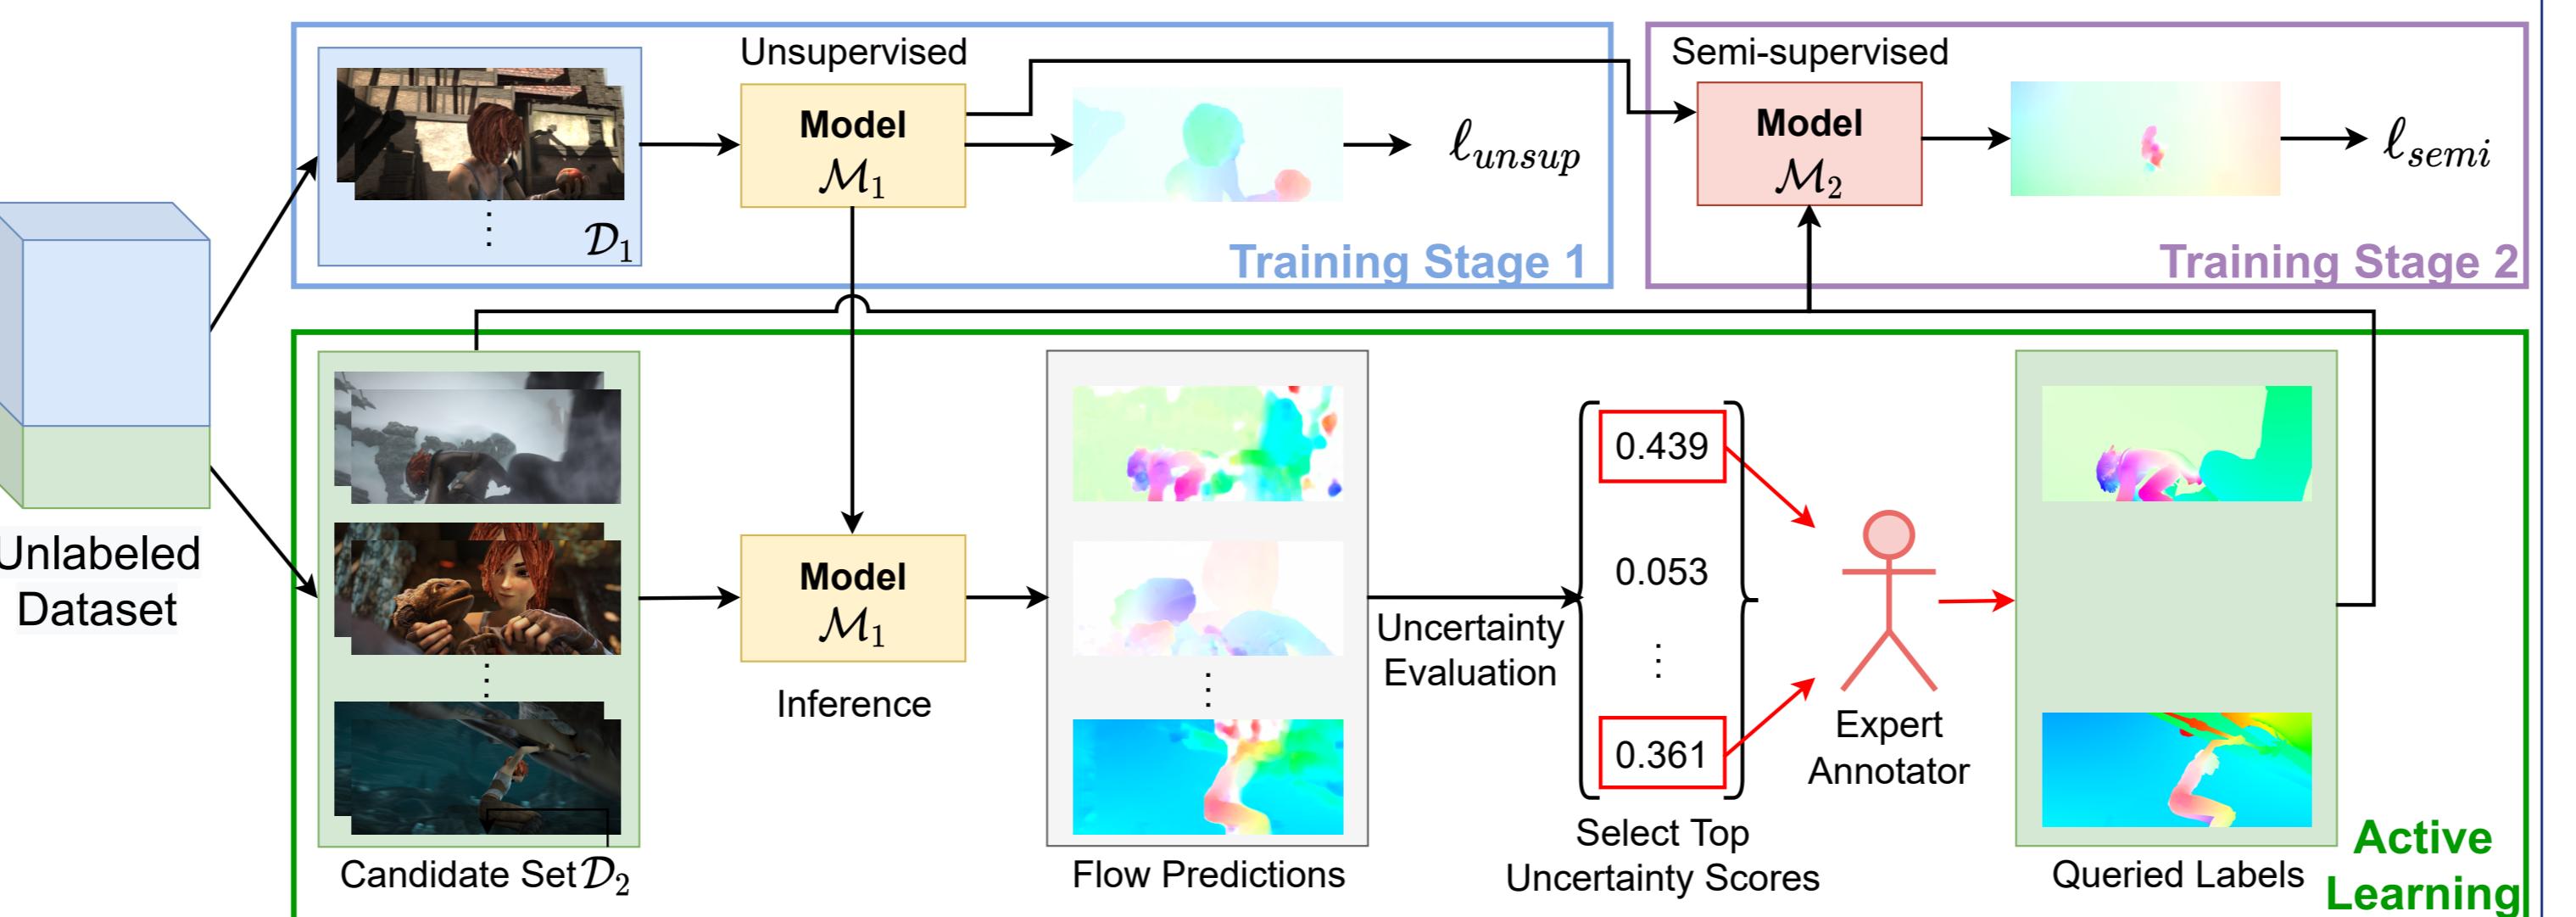
Active learning picks samples for labeling

- Given a label budget, **which** samples to label?
- Random Selection → **Active Learning**
- **New problem setup**: Semi-supervised optical flow estimation under certain label budget
- **Active learning picks most effective samples to label** → fully use label budget

Network Architecture

- **Network**: unsup SOTA: ARFlow[1]
- **Unsup loss**: multi-scale bidirectional photometric loss, smoothness loss, and augmentation loss as in ARFlow[1]
- **Sup loss**: multi-scale robust L_1 loss as in PWC-Net[2]
- **Semi-sup loss**: mixed supervised and unsupervised loss
- **Code**: <https://github.com/duke-vision/optical-flow-active-learning-release>

Active Learning Method



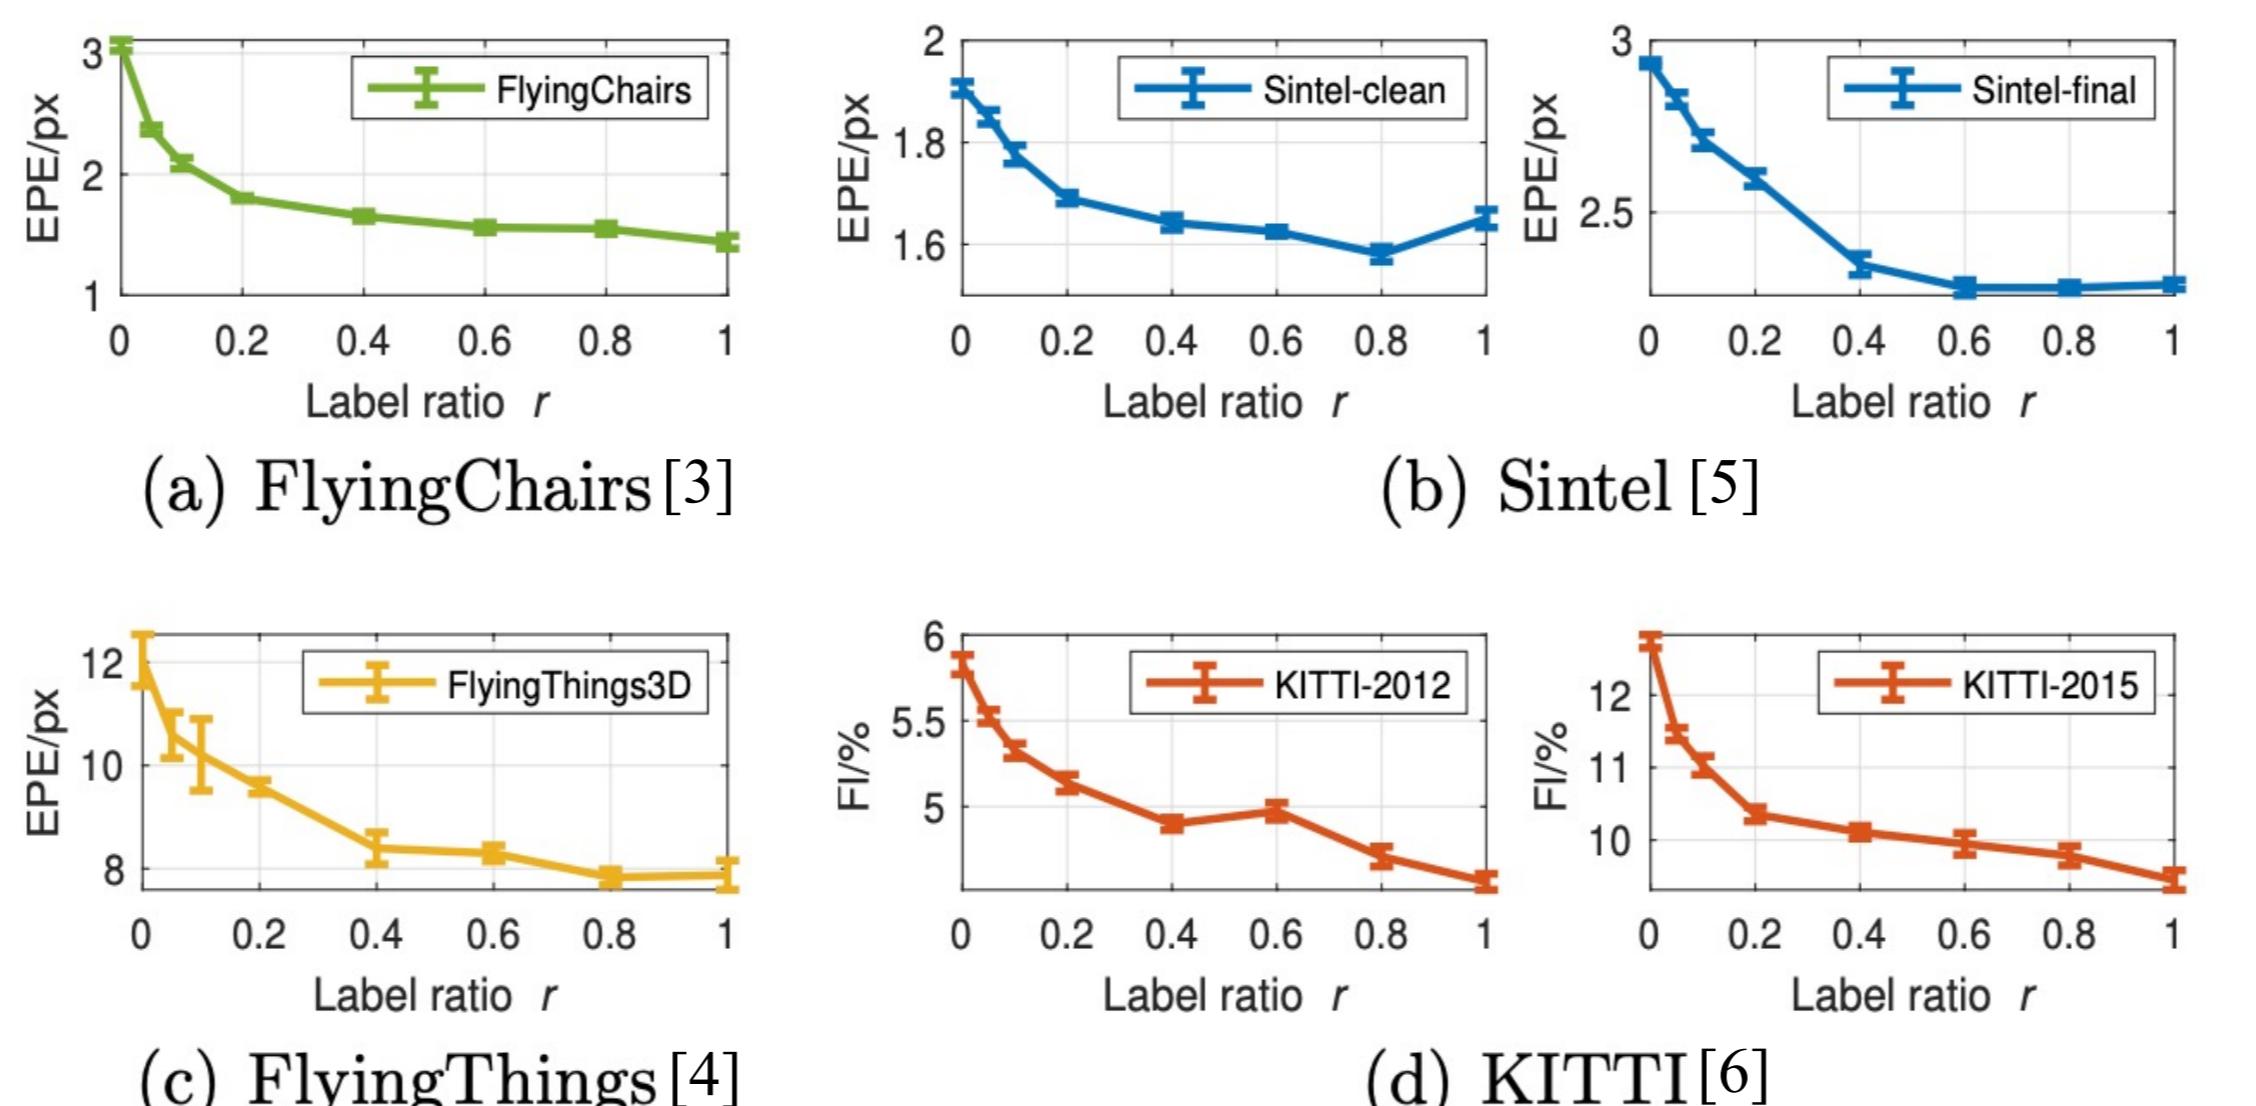
1. **Training Stage 1**: Train an unsup model \mathcal{M}_1 on \mathcal{D}_1
2. **Active Learning**: Use \mathcal{M}_1 to infer flow on \mathcal{D}_2 and evaluate uncertainty scores. Top samples query labels
3. **Training Stage 2**: Fine-tune \mathcal{M}_1 on the partially labeled \mathcal{D}_2 with semi-supervised training to get final model \mathcal{M}_2

Uncertainty scores

- *Photo loss*: photometric loss
- *Occ ratio*: ratio of occluded pixels from consistency check
- *Flow grad norm*: estimated flow’s gradient magnitude

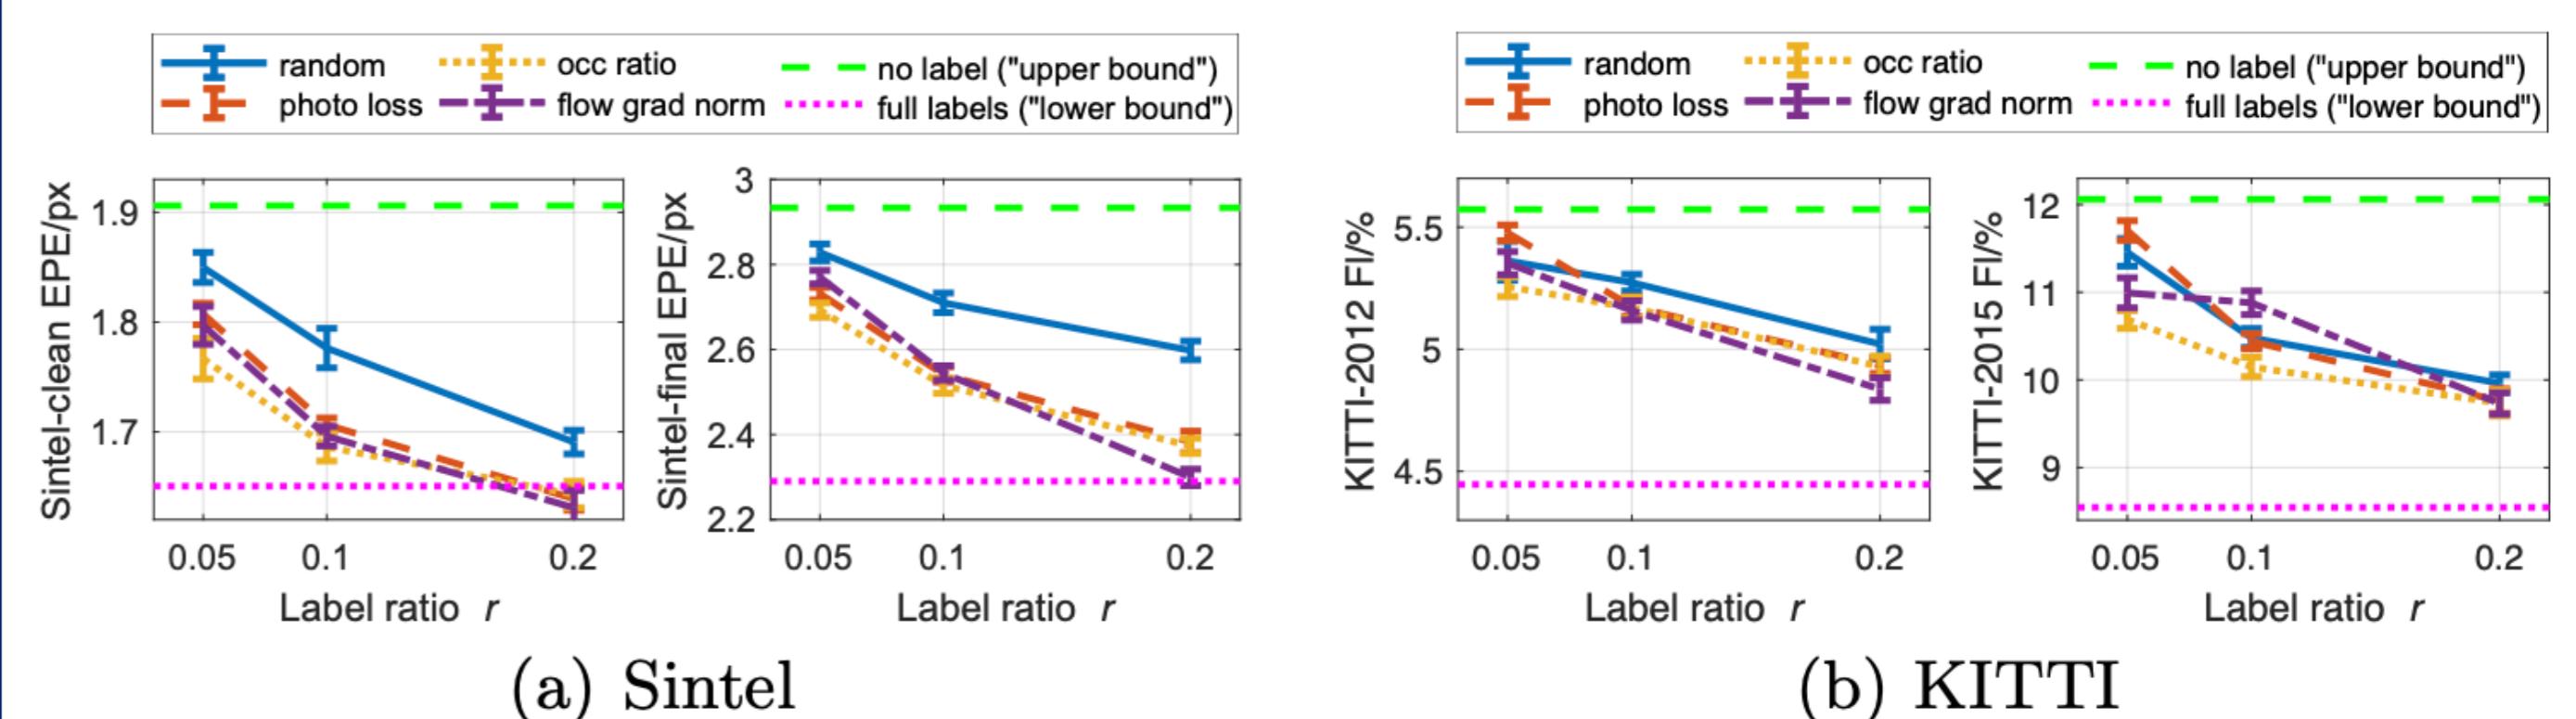
Semi-supervised Training Results

“Error-label ratio” curves: even a few labels help a lot



Active Learning Results

Validation: occ ratio works better than baseline



Benchmark test: occ ratio works (even w/ 5-20% labels)

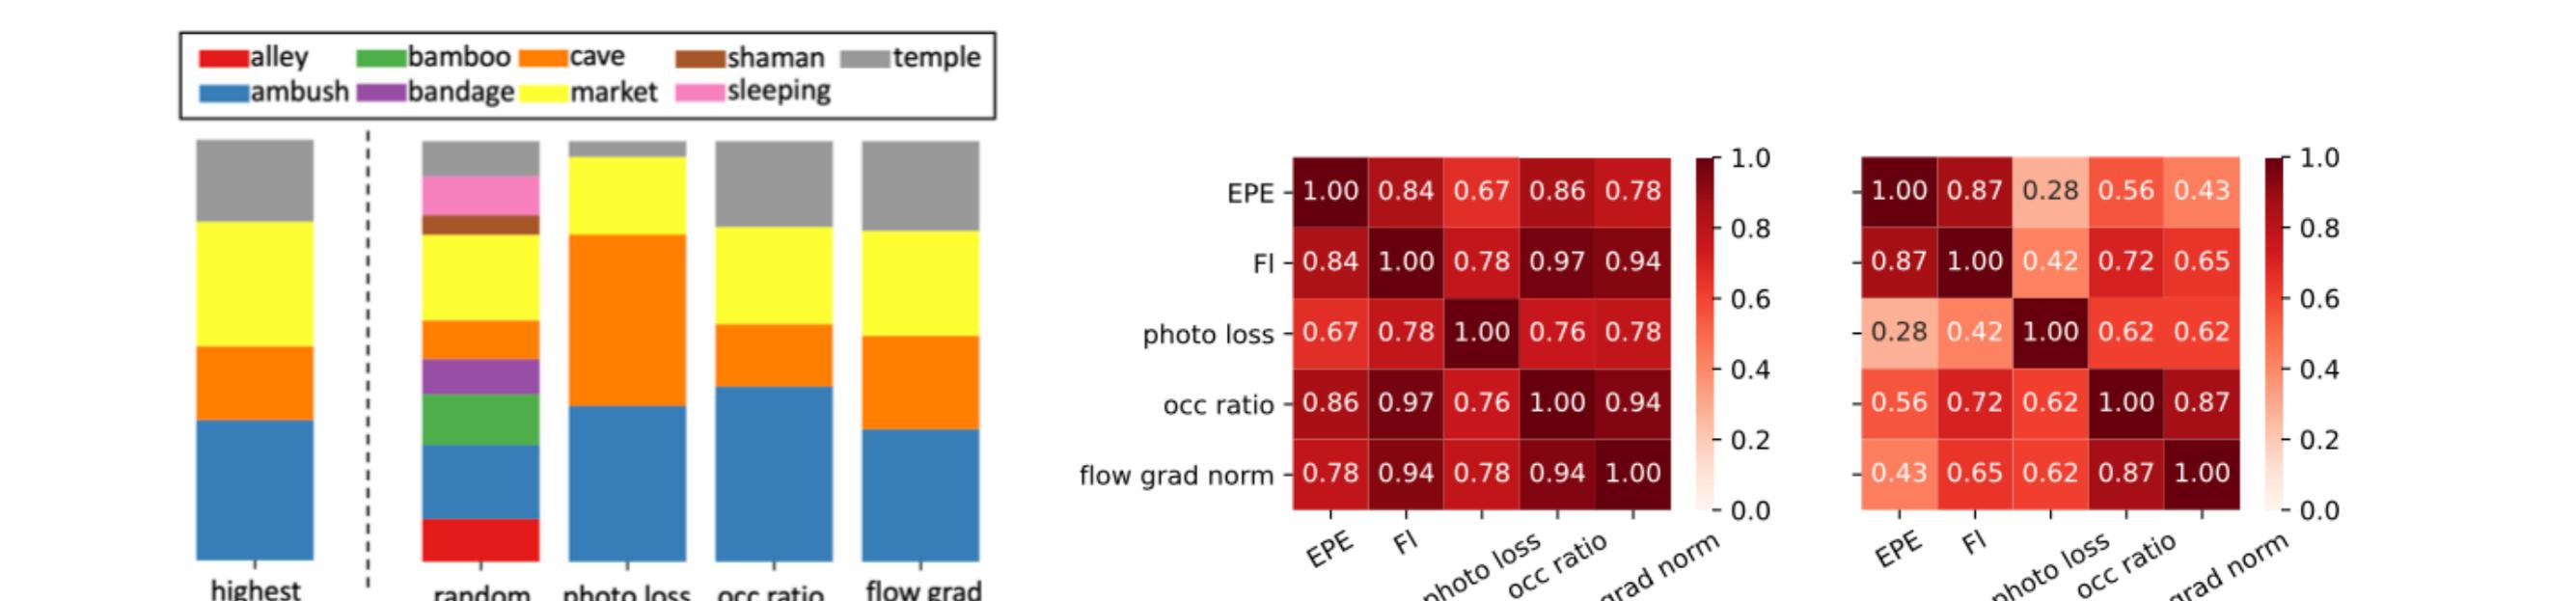
Label ratio r	Method	Train			Test		
		Clean	Final	all	Clean	Final	all
unsup							
$r = 0$	SelFlow [29]	(2.88)	(3.87)	6.56	2.67	38.30	6.57
	UFlow [19]	(2.50)	(3.39)	5.21	2.01	31.06	6.58
	ARFlow [28]	(2.79)	(3.73)	4.78	1.91	28.26	5.89
$r = 0.05$	Ours(rand)	(2.09)	(2.99)	4.04	1.52	24.65	5.49
	Ours(occ)	(1.95)	(2.38)	4.11	1.61	24.39	5.28
	Ours(occ-2x)	(1.94)	(2.55)	3.98	1.45	24.42	5.35
$r = 0.1$	Ours(rand)	(2.36)	(3.18)	3.91	1.47	23.82	5.21
	Ours(occ)	(1.64)	(1.98)	4.28	1.65	25.49	5.31
	Ours(occ-2x)	(1.75)	(2.30)	4.06	1.61	23.94	5.09
$r = 0.2$	Ours(rand)	(2.17)	(2.93)	3.89	1.56	22.86	5.20
	Ours(occ)	(1.35)	(1.63)	4.36	1.86	24.76	5.09
	Ours(occ-2x)	(1.57)	(2.05)	3.79	1.44	23.02	4.62
$r = 0.5$	PWC-Net [47]	(2.02)	(2.08)	4.39	1.72	26.17	5.04
	IRR-PWC [15]	(1.92)	(2.51)	3.84	1.47	23.22	4.58
	RAFT [49]	(0.77)	(1.27)	1.61	0.62	9.65	2.86
$r = 1$	RAFT [49]	-	-	-	-	-	-

(a) Sintel

(b) KITTI

Model Analysis

Active learning picks challenging samples to label



References

- [1] Liu, L., et al.: Learning by analogy: Reliable supervision from transformations for unsupervised optical flow estimation. In: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition. pp. 6489–6498 (2020)
- [2] Sun, D., et al.: PWC-Net: CNNs for optical flow using pyramid, warping, and cost volume. In: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition. pp. 8934–8943 (2018)
- [3] Dosovitskiy, A., et al.: Flownet: Learning optical flow with convolutional networks. In: Proceedings of the IEEE International Conference on Computer Vision. pp. 2758–2766 (2015)
- [4] Mayer, N., et al.: A large dataset to train convolutional networks for disparity, optical flow, and scene flow estimation. In: Proceedings of the IEEE International Conference on Computer Vision (2016)
- [5] Butler, D.J., et al.: A naturalistic open source movie for optical flow evaluation. In: Proceedings of the European Conference on Computer Vision. pp. 611–625. Part IV, LNCS 7577, Springer-Verlag (2012)
- [6] Menze, M. and Geiger, A.: Object scene flow for autonomous vehicles. In: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (2015)



Acknowledgments: This material is based upon work supported by the National Science Foundation under Grant No. 1909821 and by the Intelligence Advanced Research Projects Agency under contract number 2021-21040700001.